

Problem Set 1

Machine Learning AT70.9022

- Estimate a linear regression model using the normal equations

For estimating the linear regression model using the normal equation, first we need to get the dataset. The data set we are using for this model is 'weight.txt' which has 80 elements. By using height and age, weight will be predicted.

To find θ , this formula is used $\theta = (X^T X)^{-1} X^T y$

Once we have θ , we can calculate

$$J(\theta) = \frac{1}{2} \sum_{i=1}^m (h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})^2.$$

Which will give us the cost we get by using the hypothesis.

The python implementation of this experiment is given below:

```
import matplotlib.pyplot as plt
import numpy
from mpl_toolkits.mplot3d import Axes3D

#3D
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

fig.subplots_adjust(top=0.85)
ax.set_title('Normal Equation')
```

```
#importing dataset
datain = numpy.loadtxt('weight.txt')

x0 = numpy.matrix(numpy.ones(100)).transpose()
X = numpy.concatenate((x0,datain[:,0:2]),1)
Y = datain[:,2:]

Xtrain = X[:-20,:]
Xtest = X[-20:,:]
Ytrain = Y[:-20,:]
Ytest = Y[-20:,:]

# Finding the theta
theta_opt = numpy.linalg.inv(Xtrain.transpose() * Xtrain) * Xtrain.transpose() * Ytrain

# plotting actual values and predicted value
ax.scatter(Xtest[:,1],Xtest[:,2], Xtest* theta_opt)
ax.scatter(Xtest[:,1],Xtest[:,2], Ytest,c='red')

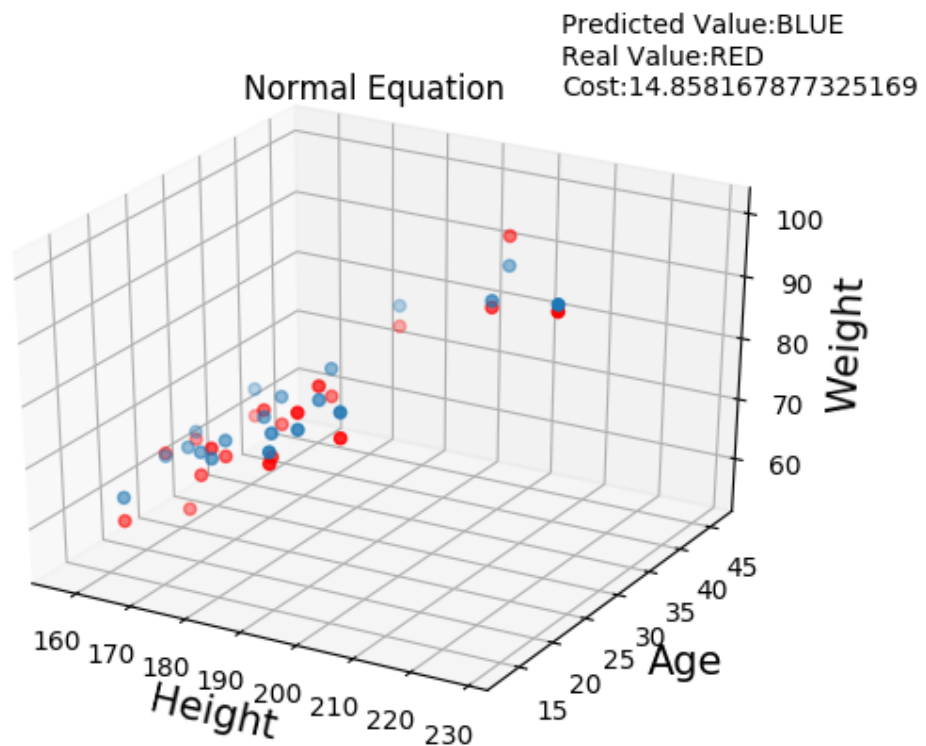
#The Cost
cost_function = ((Xtest * theta_opt - Ytest).transpose() * (Xtest * theta_opt - Ytest))/
len(Xtrain)

ax.text2D(0.70, 0.95, "Predicted Value:BLUE\nReal
Value:RED\nCost:"+str(cost_function[0,0]), transform=ax.transAxes)
ax.set_xlabel('Height', fontsize=15, rotation = 0)
ax.set_ylabel('Age', fontsize=15, rotation = 0)
ax.set_zlabel('Weight', fontsize=15, rotation = 0)
```

plt.show()

Observation:

After calculating the θ we test and see the observed and predicted values in the graph given below:



We can see that the cost we are getting is very small 14.8 approximate. This seems to be a very good cost. Therefore we can call this as a good hypothesis.

- Implement batch gradient descent.

For implementing the batch gradient descent we again take the dataset. First we take the initial theta as [0,1,1]. We use below formula to calculate the new θ :

$$\theta^{(n+1)} = \theta^{(n)} - \alpha \nabla_J(\theta^{(n)}).$$

After 10 iterations we use the θ to predict the 'weights'. And in the end we find the cost we get after using the hypothesis.

The python implementation of this experiment is given below:

```
import matplotlib.pyplot as plt

import numpy

from mpl_toolkits.mplot3d import Axes3D

#3D

fig = plt.figure()

ax = fig.add_subplot(111, projection='3d')

fig.subplots_adjust(top=0.85)

ax.set_title('Batch gradient')

# importing data

datain = numpy.loadtxt('weight.txt')

x0 = numpy.matrix(numpy.ones(100)).transpose()
```

```
X = numpy.concatenate((x0,datain[:,0:2]),1)

Y = datain[:,2:]


Xtrain = X[:-20,:]
Xtest = X[-20:,:]
Ytrain = Y[:-20,:]
Ytest = Y[-20:,:]


# initial theta values

theta_cur = numpy.matrix([0, 1, 1]).transpose()


storing_residual = numpy.empty([10, 1])


for iter in range(0,10):

    # calculating the residual

    residual = (Xtrain * theta_cur - Ytrain).transpose() * (Xtrain * theta_cur - Ytrain)


    # recording

    storing_residual[iter,0] = residual


    # delta values

    delta_intercept = numpy.sum(Xtrain * theta_cur - Ytrain)

    delta_height = numpy.sum((Xtrain * theta_cur - Ytrain).transpose()*Xtrain[:,1])
```

```
delta_age = numpy.sum((Xtrain * theta_cur - Ytrain).transpose()*Xtrain[:,2])

# gradient

gradient = numpy.matrix([delta_intercept, delta_height, delta_age]).transpose()

# new theta value

theta_cur = theta_cur - 0.00001/len(Xtrain) * gradient

cost_function = ((Xtest * theta_cur - Ytest).transpose() * (Xtest * theta_cur - Ytest))/
len(Xtest)

# Plotting 3d graph for actual values and predicted values

ax.scatter(Xtest[:,1],Xtest[:,2], Xtest* theta_cur)

ax.scatter(Xtest[:,1],Xtest[:,2], Ytest,c='red')

ax.text2D(0.70, 0.95, "Predicted Value:BLUE\nReal
Value:RED\nCost:"+str(cost_function[0,0]), transform=ax.transAxes)

ax.set_xlabel('Height', fontsize=15, rotation = 0)

ax.set_ylabel('Age', fontsize=15, rotation = 0)

ax.set_zlabel('Weight', fontsize=15, rotation = 0)

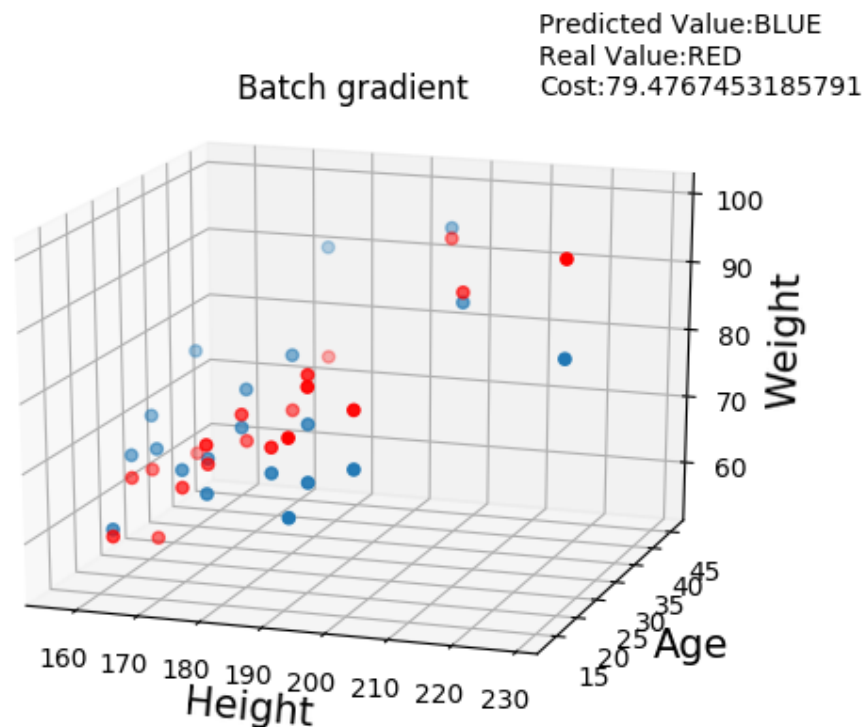
plt.show()

plt.title('epochs vs.  $J(\theta)$ ')
```


```
plt.plot(range(0,10), storing_residual/len(Xtrain))  
  
plt.scatter(range(0,10), storing_residual/len(Xtrain))  
  
plt.xlabel('epochs')  
  
plt.ylabel('Cost')  
  
plt.show()
```

Observation:

After calculating the θ we test and see the observed and predicted values in the graph given below:





We can see that the cost doesn't seem to be as good as the Normal equation ones. It might get better if we play around with the initial value and the learning rate.



The cost seems to go down really fast in just 10 iterations. But not close to the cost as in normal equation.

After playing around with the initial values of θ , these are the significant changes that we can see.



- Implement stochastic gradient descent.

For implementing the stochastic gradient descent we again take the dataset. First we take the initial theta as $[0,1,1]$. We use below formula to calculate the new θ :

```

 $\theta \leftarrow$  some initial guess
While not converged
  For  $i \in 1..m$ 
     $\theta \leftarrow \theta - \alpha(h_{\theta}(\mathbf{x}^{(i)}) - y^{(i)})\mathbf{x}^{(i)}$ 

```

Here rather than calculating the gradient with the whole dataset we just take one elements from the dataset and then find the gradient and get the new θ value. In the end we will find the cost using our test values.

The python implementation of this experiment is given below:

```

import matplotlib.pyplot as plt
import numpy
from mpl_toolkits.mplot3d import Axes3D

# 3D
fig = plt.figure()
ax = fig.add_subplot(111, projection='3d')

fig.subplots_adjust(top=0.85)
ax.set_title('Stochastic Gradient Descent')

# Importing data
datain = numpy.loadtxt('weight.txt')

```

```

x0 = numpy.matrix(numpy.ones(100)).transpose()
X = numpy.concatenate((x0,datain[:,0:2]),1)
Y = datain[:,2:]

Xtrain = X[:-20,:]
Xtest = X[-20:,:]
Ytrain = Y[:-20,:]
Ytest = Y[-20:,:]

# initial values for thetas
theta_cur = numpy.matrix([-28.0, 0.5, 0.2]).transpose()

# Storing squared residuals
storing_residual = numpy.empty([10, 1])

for iter in range(0,10):
    # getting the squared residuals
    residual = (Xtrain * theta_cur - Ytrain).transpose() * (Xtrain * theta_cur - Ytrain)

    # recording
    storing_residual[iter,0] = residual

# iteration to find gradient after every single dataset element
for x in range(0,len(Xtrain)):
    delta_intercept = (Xtrain[x,:] * theta_cur - Ytrain[x])[0,0]
    delta_height = ((Xtrain[x,:] * theta_cur - Ytrain[x]) * Xtrain[x,1])[0,0]
    delta_age = ((Xtrain[x,:] * theta_cur - Ytrain[x]) * Xtrain[x,2])[0,0]

```

```
gradient = numpy.matrix([delta_intercept, delta_height, delta_age]).transpose()
theta_cur = theta_cur - 0.00001/len(Xtrain) * gradient

cost_function = ((Xtest * theta_cur - Ytest).transpose() * (Xtest * theta_cur - Ytest))/
len(Xtest)

# Plotting 3d graph for actual values and predicted values
ax.scatter(Xtest[:,1],Xtest[:,2], Xtest* theta_cur)
ax.scatter(Xtest[:,1],Xtest[:,2], Ytest,c='red')
ax.text2D(0.70, 1, "Predicted Value:BLUE\nReal
Value:RED\nCost:"+str(cost_function[0,0]), transform=ax.transAxes)
ax.set_xlabel('Height', fontsize=15, rotation = 0)
ax.set_ylabel('Age', fontsize=15, rotation = 0)
ax.set_zlabel('Weight', fontsize=15, rotation = 0)

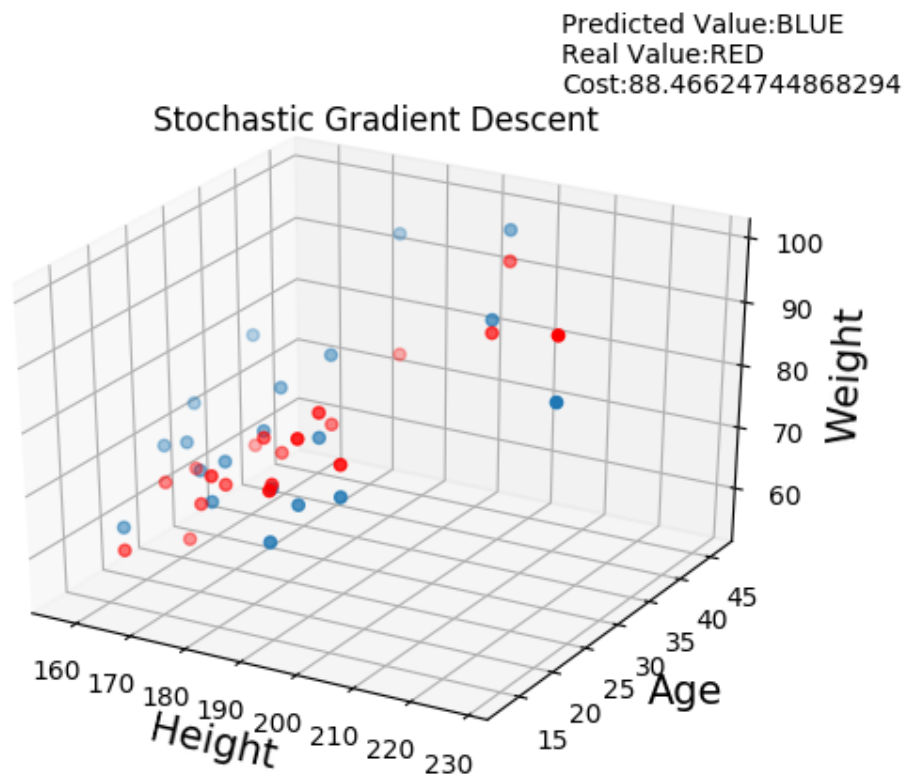
plt.show()

plt.title('epochs vs.  $J(\theta)$ ')


plt.plot(range(0,10), storing_residual/len(Xtrain))
plt.scatter(range(0,10), storing_residual/len(Xtrain))
plt.xlabel('epochs')
plt.ylabel('Cost')
plt.show()
```

Observation:

After calculating the θ we test and see the observed and predicted values in the graph given below:

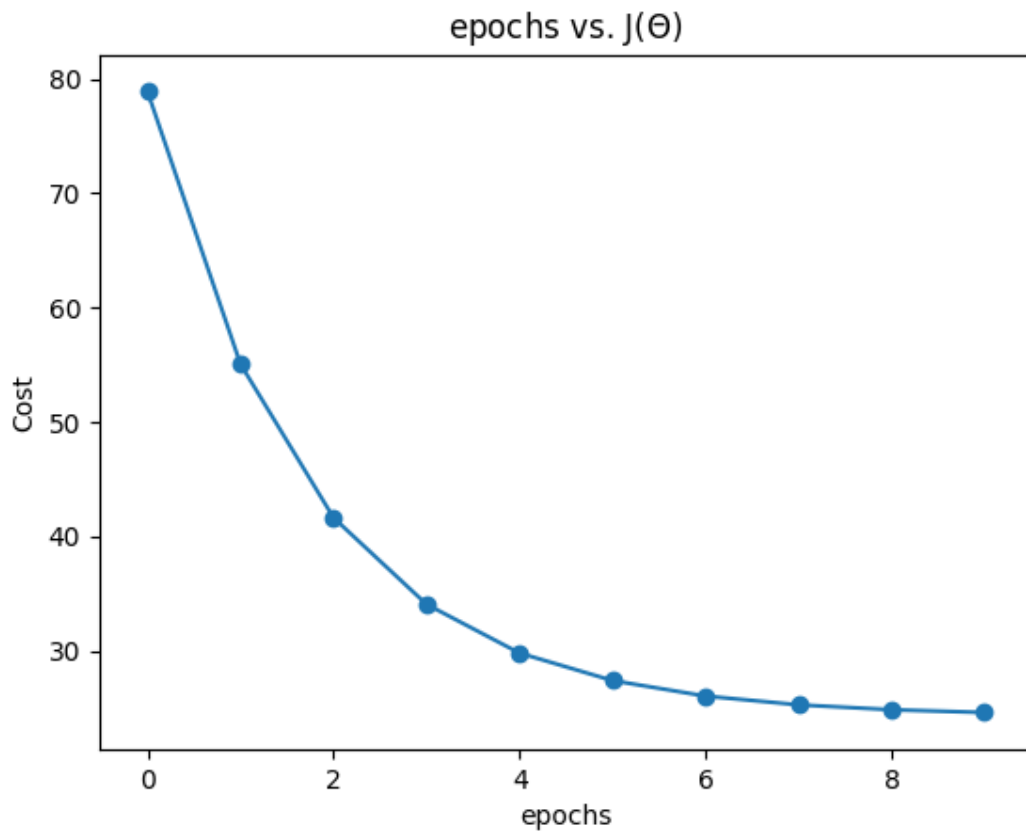



We can see that the cost isn't even better than the batch gradient stochastic. It might get better if we play around with the initial value and the learning rate.



The graph shows that the cost comes down fast but maybe it started oscillating around the minimum because of which we did not get a better cost.

After playing around with the initial values of θ , these are the significant changes that we can see.



Thursday, 18 January 2018